Technical Interview Experience - DevOps Engineer at GlobalLogic (Round 1)

Azure DevOps & CI/CD Workflows

How do you dynamically pass secrets and tokens across pipeline stages securely in Azure DevOps?

To pass secrets and tokens securely across pipeline stages in Azure DevOps, it's crucial to avoid hardcoding them directly in YAML or other pipeline definitions. Instead, utilize Azure Key Vault for storing secrets and configure pipelines to fetch them during runtime. Azure DevOps also provides mechanisms for securely handling variables and access tokens, which should be leveraged to prevent exposure of sensitive information.

Here's a breakdown of recommended practices:

- 1. Store Secrets in Azure Key Vault:
 - Azure Key Vault is the recommended service for storing secrets securely.
 - It provides role-based access control, allowing you to restrict access to specific secrets or groups of secrets.
 - You can manage secrets, keys, and certificates within Key Vault.
- 2. Access Key Vault Secrets in Pipelines:
 - Azure CLI Task: Use the AzureCLI@2 task to execute commands within your pipeline to retrieve secrets from Key Vault.
 - **Service Connection:** Establish a service connection to your Azure subscription, which will be used by the Azure CLI task to authenticate with Key Vault.
 - **Key Vault Name:** Specify the name of your Key Vault instance in the AzureCLI@2 task.
 - Secret Name/ID: Retrieve the specific secret you need using the Azure CLI commands.
- 3. Secure Variable Handling:
 - Secret Variables:

Mark variables as "secret" in your pipeline settings or variable groups to prevent them from being displayed in logs or stored in plain text.

Variable Groups:

Utilize variable groups to manage and share secrets across multiple pipelines.

• Environment Variables:

Map secrets to environment variables within tasks to make them available to scripts.

Avoid Command Line:

Never pass secrets directly on the command line.

4. Secure Access Tokens:

Job Access Tokens:

Azure Pipelines automatically generate job access tokens for each job. These tokens grant access to Azure DevOps resources.

Token Permissions:

Configure the scope of the job access token to limit the resources it can access.

Federated Identity:

Use Federated Identity for dynamic authentication with Azure resources. This eliminates the need to store long-lived credentials.

5. YAML Pipelines:

• YAML over Classic:

Use YAML pipelines for more flexibility and control over security configurations.

Templates:

Leverage templates to enforce security best practices across your pipelines.

Example (Illustrative):

Code

Example using Azure CLI and Key Vault

- task: AzureCLI@2

displayName: 'Get Secret from Key Vault'

inputs:

azureSubscription: 'your-service-connection-name'

scriptType: 'bash'

.: - .- . !:.- !:.- - **0** - ..:.

scriptLocation: 'inlineScript'

inlineScript: |

secretValue=\$(az keyvault secret show --name "your-secret-name" --vault-name "your-key-vault-name" --query "value" --output tsv)

echo "##vso[task.setvariable variable=mySecret;issecret=true]\$secretValue"

addSpnToEnvironment: true

- task: PowerShell@2

displayName: 'Use the Secret'

inputs:

targetType: 'inline'

script: |

Write-Host "Using secret: \$(\$env:mySecret)" # Access the secret via environment

variable

Key Considerations:

Principle of Least Privilege:

Grant the minimum necessary permissions to your pipelines and service accounts.

Regular Audits:

Periodically review your pipeline configurations and security settings.

Stay Updated:

Keep your Azure DevOps platform and tools up to date with the latest security patches.

Monitor Activity:

Utilize Azure Monitor to track and monitor pipeline activity, including access to secrets and resources.

What are environment checks and how do they help in managing gated deployments?

Environment checks play a crucial role in managing gated deployments by ensuring that deployments only proceed when specific criteria are met. These checks act as safeguards, verifying the readiness and quality of a deployment before it moves to the next stage, preventing potential issues from reaching production.

Here's how environment checks help in managing gated deployments:

1. Controlling Deployment Progression:

• Manual Approvals:

Environments can be configured to require manual approvals before deployments can proceed. This allows designated users to review changes and assess their readiness for a specific environment, like production, before granting permission.

Automated Checks:

Automated checks can be implemented to evaluate various aspects of the deployment, such as code quality, security vulnerabilities, performance metrics, and resource availability. These checks can be run before or after deployments to ensure that the environment meets predefined standards.

Branching and Access Control:

Environments can be configured to restrict deployments from certain branches or limit access to specific resources. This helps in managing the flow of changes and ensures that only authorized code is deployed to specific environments.

2. Ensuring Deployment Quality:

Preventing Rollbacks:

By validating deployments through environment checks, the likelihood of encountering issues after deployment is reduced, minimizing the need for rollbacks and associated downtime.

Early Issue Detection:

Automated checks can identify potential problems early in the deployment process, allowing for prompt remediation before they escalate into major issues.

Increased Stability:

Rigorous environment checks contribute to more stable and reliable deployments, leading to a better user experience and reduced operational overhead.

3. Enhancing Security:

Secrets Management:

Environments can securely store secrets (like API keys or passwords) and only grant access to them after the necessary checks have passed.

Compliance:

Environment checks can be used to enforce compliance with security policies and regulations, ensuring that deployments adhere to predefined standards.

4. Facilitating Continuous Improvement:

Feedback Loop:

Environment checks can provide valuable feedback on the deployment process, allowing teams to identify areas for improvement and optimize their deployment pipeline.

Data-Driven Decisions:

Monitoring the results of environment checks can provide data-driven insights into the health and performance of deployments, enabling teams to make informed decisions about future deployments.

In essence, environment checks act as a critical component of gated deployments, providing a robust framework for ensuring that changes are validated, secure, and stable before they are released to users. They help teams to confidently and efficiently manage deployments, minimizing risks and maximizing the benefits of a continuous delivery approach.

How would you handle rollback in a multi-stage YAML pipeline if stage 3 fails but stage 1 & 2 succeeded?

To handle rollbacks in a multi-stage pipeline when a failure occurs in a later stage (e.g., stage 3) after earlier stages (e.g., stages 1 & 2) have succeeded, you need to implement a rollback strategy. This strategy should involve: 1) detecting the failure in stage 3, 2) triggering a rollback process, and 3) reverting the changes made by the successful stages (1 & 2) or at least ensuring a consistent state.

Here's a breakdown of how to achieve this, along with considerations for different pipeline systems:

1. Detecting Failure and Triggering Rollback:

• Pipeline Configuration:

Most CI/CD systems (like Azure Pipelines, AWS CodePipeline, Jenkins, etc.) allow you to define failure conditions and trigger actions based on those conditions.

Failure Strategy:

Configure the pipeline to detect a failure in stage 3. This often involves setting a failure condition (e.g., "on error" or "if stage fails") and specifying the action to be taken (e.g., "rollback" or "run a specific task").

Rollback Stage/Job:

Create a dedicated stage or job (or use a pre-defined rollback action) that will be executed when the failure condition is met.

2. Reverting Changes (Rollback):

Automated Rollbacks:

Some systems (like AWS CodePipeline) offer automatic rollback capabilities. When enabled, they automatically revert to the most recent successful execution of the failed stage.

Scripted Rollbacks:

For more complex scenarios or when automatic rollback isn't sufficient, you'll need to script the rollback process.

Database Migrations:

If database changes were involved, ensure your rollback script handles reverting these changes. This might involve rolling back database migrations or restoring from a backup.

Configuration Management:

If configuration changes were made, your rollback script should handle reverting to the previous configuration. This might involve using a configuration management tool or applying specific commands.

Application Deployments:

If deployments were involved, your rollback script might need to redeploy a previous version of your application.

• Idempotency:

Make sure your rollback scripts are idempotent, meaning they can be run multiple times without causing unintended side effects. This is crucial for handling potential issues during the rollback process.

3. Examples in Different Systems:

• Azure Pipelines:

Azure Pipelines allows you to define stages with dependencies. You can define stage B to run only if stage A fails and set up stage C to run only if stage B succeeds. You can also use custom conditions to control when stages are executed, allowing for rollback scenarios.

AWS CodePipeline:

AWS CodePipeline supports automatic rollback. You can enable automatic rollback on stage failure, and it will revert to the most recent successful execution.

Jenkins:

Jenkins can be configured to run specific scripts or jobs on failure. You can use || in shell scripts to execute a rollback command if the previous command fails.

4. Considerations:

Version Control:

Ensure that your rollback process is integrated with version control. This will allow you to easily revert to previous versions of code and configurations.

Testing:

Thoroughly test your rollback strategy to ensure it works as expected and that you can recover from failures.

Documentation:

Maintain clear documentation of your rollback process, including the steps involved and the scripts used.

Monitoring:

Implement monitoring to detect failures quickly and trigger rollbacks promptly.

In summary, a robust rollback strategy involves detecting failures, triggering a rollback process, and reverting changes in a controlled and repeatable manner, with consideration for the specific tools and environment being used.

✓ How do you implement reusability across pipelines for 15+ microservices?

To implement reusability across pipelines for 15+ microservices, focus on creating modular, reusable pipeline components and leveraging tools that support pipeline-ascode. This involves defining common pipeline stages, using configuration management for service-specific settings, and employing tools like Tekton or GitLab CI/CD to manage and automate pipeline execution.

Here's a breakdown of how to achieve reusability:

1. Define Reusable Pipeline Stages:

Modularize Pipeline Steps:

Break down the pipeline into smaller, reusable stages like "Build," "Test," "Deploy," and "Security Scan".

Parameterize for Flexibility:

Use parameters within these stages to handle variations between microservices, such as different programming languages, dependencies, or deployment targets.

Centralized Configuration:

Store service-specific configurations (e.g., environment variables, connection strings) in a central location (e.g., a configuration management system) and reference them in the pipeline.

Version Control for Pipelines:

Treat pipeline definitions as code (e.g., using YAML files) and store them in a version control system like Git. This enables branching, versioning, and collaboration on pipelines.

2. Leverage CI/CD Tools:

Choose Appropriate Tools:

Select CI/CD tools that support pipeline-as-code and offer features like reusable pipelines and task management. Options include Tekton, Jenkins, GitLab CI/CD, and others.

Implement Pipeline as Code:

Define your pipelines using a declarative approach (e.g., YAML) that allows for versioning, reusability, and easier management.

Automate Pipeline Execution:

Configure triggers (e.g., Git commits, pull requests) to automatically trigger pipeline execution when code changes are made.

3. Manage Dependencies and Artifacts:

Containerization:

Containerize each microservice using tools like Docker to ensure consistency across different environments.

• Artifact Repositories:

Utilize artifact repositories (e.g., Docker Hub, Nexus) to store and manage built images and other artifacts.

• Service Discovery:

Implement service discovery mechanisms (e.g., using Kubernetes or Consul) to allow microservices to locate and communicate with each other.

4. Promote Collaboration and Knowledge Sharing:

Documentation:

Maintain clear and up-to-date documentation for each microservice and its associated pipeline.

Standardization:

Establish common standards and best practices for microservice development and deployment.

Shared Libraries:

Consider creating shared libraries for common functionality (e.g., logging, authentication) that can be reused across multiple microservices.

5. Consider Advanced Techniques:

Canary Deployments:

Implement canary deployments to test new versions of a microservice with a small subset of users before rolling it out to the entire user base.

• Blue/Green Deployments:

Use blue/green deployments to switch between different versions of a microservice with minimal downtime.

Feature Flags:

Utilize feature flags to enable or disable specific features within a microservice, allowing for more controlled releases and experimentation.

By adopting these strategies, you can create a robust and reusable CI/CD pipeline infrastructure that effectively supports the development and deployment of your 15+ microservices.

Kubernetes & Container Platform

How do you manage Kubernetes secrets securely without exposing them in Git or CI logs?

To manage Kubernetes secrets securely without exposing them in Git or CI logs, leverage external secret management tools, encrypt secrets at rest, and control access using RBAC. Consider tools like HashiCorp Vault, AWS Secrets Manager, or Azure Key Vault, and encrypt secrets before storing them in etcd. Additionally, use Sealed Secrets for encrypting secrets for GitOps workflows.

Here's a more detailed breakdown:

1. External Secret Management:

Use dedicated tools:

Instead of relying solely on Kubernetes Secrets, integrate with external secret management systems like <u>HashiCorp Vault</u>, <u>AWS Secrets Manager</u>, or Azure Key Vault. These tools offer advanced features like encryption, access control, and auditing.

Benefits:

Centralized storage, granular access control, audit trails, and rotation policies.

Integration:

Tools like the Secrets Store CSI Driver and External Secrets Operator enable seamless integration with Kubernetes.

2. Encryption at Rest:

• Enable encryption:

Configure Kubernetes to encrypt secrets at rest within etcd using a KMS provider. This ensures that even if etcd is compromised, the secrets remain encrypted.

• Dual Envelope Encryption:

Combine Kubernetes encryption at rest with client-side encryption (e.g., using tools like Mozilla SOPS) or KMS for an extra layer of security.

3. GitOps and Sealed Secrets:

• Sealed Secrets: Use Sealed Secrets to encrypt secrets before committing them to Git. The Sealed Secrets controller in the cluster decrypts them at runtime.

- **Client-side encryption:** Tools like Mozilla SOPS allow you to encrypt secrets on your local machine before committing them to Git.
- Avoid storing plaintext secrets: Never store raw secrets in your Git repositories or CI/CD logs.

4. Access Control (RBAC):

Least Privilege:

Implement Role-Based Access Control (RBAC) to restrict access to secrets based on roles and permissions. Grant only the minimum necessary access to specific users, service accounts, and pods.

Namespace isolation:

Utilize Kubernetes namespaces to further isolate secrets and limit their scope.

5. Other Best Practices:

- Regular Rotation: Rotate secrets regularly to minimize the impact of potential leaks.
- Audit Logging: Enable audit logging to track secret access and identify potential security breaches.
- **Secret Scanning:** Use tools to scan your code and configuration files for accidentally exposed secrets.
- Monitoring: Monitor secret usage to detect suspicious activity.
- **Environment Variables:** Use environment variables sparingly for secrets and prefer mounting secrets as volumes.
- **Secure Communication:** Ensure that communication between your applications and the Kubernetes API server is encrypted using TLS.

By implementing these practices, you can significantly enhance the security of your Kubernetes secrets and protect sensitive information from unauthorized access and exposure.

A new deployment shows no logs and isn't reachable — walk me through your end-to-end debug strategy.

A systematic approach to debugging a failed deployment that's unreachable and shows no logs involves checking deployment status, infrastructure, network configurations, application logs (if available), and service configurations. Start by verifying the deployment history and logs, then examine the underlying infrastructure and network settings. Finally, inspect application logs and relevant service configurations.

1. Verify Deployment Status:

- **Check Deployment History:** Examine the deployment history for the application. This will often provide error messages or indicators of failure.
- **Correlation IDs:** If available, note any correlation IDs associated with the deployment. These can help track down the specific issues.
- **Deployment Status:** Determine if the deployment is in a failed, pending, or completed state.

2. Investigate Infrastructure:

- **Resource Availability:** Ensure sufficient resources (CPU, memory, storage) are available on the host or cluster where the application is deployed.
- Resource Limits: Verify that resource limits haven't been exceeded.
- Hardware Issues: Check for any hardware-related problems on the underlying servers.

3. Examine Network Configurations:

- Firewall Rules: Ensure that firewall rules are not blocking traffic to the application.
- Load Balancer: If using a load balancer, verify its configuration and health checks.
- **DNS Resolution:** Confirm that the application's domain name resolves to the correct IP address.
- **Network Connectivity:** Verify basic network connectivity to the deployment environment.

4. Analyze Application Logs (if available):

- Log Aggregation: If logs are aggregated, check the logs for the application instance or container.
- Container Logs: If using containers, inspect the container logs for errors or issues.

• **Application-Specific Logs:** Examine any application-specific log files or output streams.

5. Review Service Configurations:

- **Service Definition:** Ensure that the service definition (e.g., Kubernetes service, Docker Compose file) is correct.
- **Service Discovery:** Verify that the service is properly registered with any service discovery mechanisms.
- **Configuration Files:** Check for any misconfigurations in the application's configuration files.

6. Consider Rollback (If Necessary):

- Rollback to Previous Version: If available, rollback to a previous working version of the application.
- Feature Flags: If feature flags are used, disable the problematic feature.

7. Test and Verify:

Basic Functionality:

After addressing the issues, test basic functionality to ensure the application is working.

End-to-End Testing:

Perform end-to-end testing to verify all aspects of the application are functioning correctly.

By following this structured approach, you can effectively debug a failed deployment and bring your application back online.

What is the difference between Init containers and sidecars — and where have you used both in real projects?

Init containers and sidecar containers are both types of containers used within a Kubernetes pod to extend functionality, but they differ in their lifecycle and purpose. Init containers run to completion before the main application containers start, while sidecar containers run concurrently with the main application containers and remain active throughout the pod's lifecycle.

Init Containers:

• Purpose:

Init containers are used to perform setup tasks that need to be completed before the main application containers can start.

• Lifecycle:

They run sequentially and must complete successfully before the main containers are started.

- Examples:
 - Downloading configuration files from a remote source.
 - Initializing a database (schema creation, data migration).
 - Setting up shared volumes.
 - Waiting for specific services or resources to become available.

Sidecar Containers:

Purpose:

Sidecar containers provide supplementary functionality or services to the main application container, often enhancing or extending its capabilities.

Lifecycle:

They run concurrently with the main application container and remain active throughout the pod's lifecycle.

- Examples:
 - Logging: Collecting and forwarding logs from the main application.
 - Monitoring: Collecting metrics from the main application.
 - Security: Handling security-related tasks like authentication or authorization.
 - Proxying: Acting as a reverse proxy for the main application.

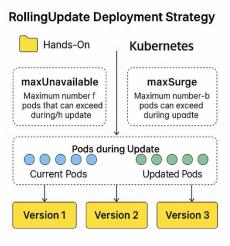
Key Differences Summarized:

Feature Init Container Sidecar Container

Purpose	Initialization tasks before main containers	Auxiliary services alongside main containers
Lifecycle	Runs to completion before main containers start	Runs concurrently with main containers and stays active
Concurrency	Sequential execution	Concurrent execution
Probes	Not supported	Supported
Restart Policy	Part of pod's restart policy	Independent restart policy

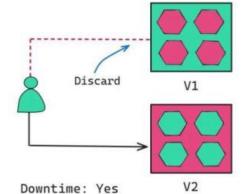
✓ How do you perform rolling updates with zero downtime in AKS?

In Azure Kubernetes Service (AKS), you can perform rolling updates with zero downtime by using the rolling update strategy in your deployments. This strategy ensures that new versions of your application are deployed incrementally, with old pods being replaced by new ones one at a time. This prevents any interruption to your application's availability during the update process.

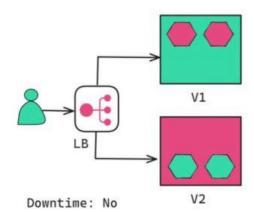




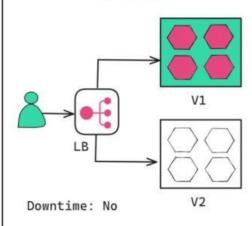
1. Recreate



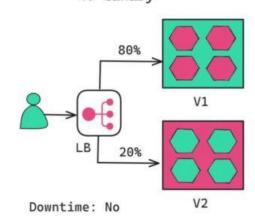
2. Rolling Update



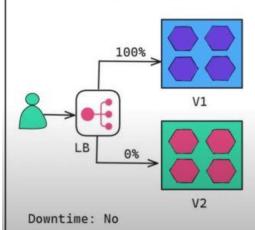
3. Shadow



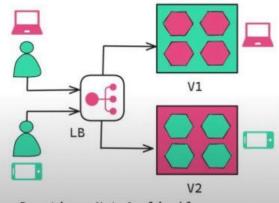
4. Canary



5. Blue-Green



6. A/B Testing



Downtime: Not Applicable

Here's a more detailed explanation:

Key Concepts:

Rolling Updates:

Kubernetes updates deployments by incrementally replacing old pods with new ones. This happens in a controlled manner, ensuring that at least a certain number of pods are always available during the update process.

maxSurge and maxUnavailable:

These settings in the deployment manifest control the rolling update behavior.

- maxSurge: Defines the maximum number of extra pods that can be created beyond the desired number of pods during the update.
- maxUnavailable: Defines the maximum number of pods that can be unavailable during the update.

• Readiness Probes:

Kubernetes uses readiness probes to check if a pod is ready to serve traffic before routing traffic to it. This ensures that only healthy pods receive traffic during the update.

Health Checks:

Configuring health checks, specifically readiness probes, is crucial for zero-downtime updates. Kubernetes uses these probes to determine if a pod is healthy and ready to accept traffic.

How it works:

1. 1. Deployment Configuration:

The deployment manifest includes the rollingUpdate strategy under spec.strategy.type.

2. 2. Incremental Update:

Kubernetes starts by creating new pods based on the updated image.

3. 3. Readiness Check:

The system waits for the new pods to become ready (as indicated by the readiness probe) before proceeding.

4. 4. Gradual Replacement:

Once a new pod is ready, an old pod is terminated, and the process repeats until all old pods are replaced.

5. 5. Rollback:

If issues are detected during the update (e.g., a readiness probe fails), the update can be rolled back to the previous stable version.

Example Deployment Manifest:

```
Code
apiVersion: apps/v1
kind: Deployment
metadata:
name: my-deployment
spec:
replicas: 3
strategy:
 type: RollingUpdate
 rollingUpdate:
  maxSurge: 1
  maxUnavailable: 0
selector:
 matchLabels:
  app: my-app
template:
 metadata:
  labels:
   app: my-app
 spec:
  containers:
  - name: my-container
   image: my-image:latest
   ports:
   - containerPort: 80
   readinessProbe:
    httpGet:
     path: /health
```

port: 80

initialDelaySeconds: 5 periodSeconds: 10

Key Points for Zero Downtime:

Sufficient Replicas:

Deploy your application with a sufficient number of replicas (at least two for production workloads) to ensure availability during updates.

Health Checks:

Implement and configure readiness probes correctly to ensure that new pods are ready to serve traffic before they are added to the service.

Gradual Updates:

Avoid abrupt changes by using rolling updates with appropriate maxSurge and maxUnavailable settings.

Monitoring and Rollback:

Monitor the update process and be prepared to roll back to a previous version if issues are detected.

By following these guidelines, you can perform rolling updates in AKS with minimal to no downtime, ensuring a smooth and continuous user experience.

Infrastructure as Code (Terraform / Bicep)

How do you structure Terraform for large teams managing different environments (dev/stage/prod)?

To manage different environments (dev/stage/prod) with Terraform in large teams, adopt a modular structure, use workspaces or separate configurations for each environment, and leverage remote state storage. This approach ensures consistency, scalability, and controlled access across environments.

Here's a breakdown of the key strategies:

1. Modularization:

• Create reusable modules:

Encapsulate infrastructure components (e.g., networking, databases, compute) into reusable modules. This promotes code reuse, reduces redundancy, and simplifies maintenance.

Organize modules logically:

Structure modules based on functionality or infrastructure components (e.g., modules/network, modules/compute, modules/database).

2. Environment Separation:

• Workspaces (CLI or Terraform Cloud/Enterprise):

Use Terraform Workspaces to manage separate state files for each environment within the same configuration. This allows you to apply changes to specific environments without affecting others.

Separate configurations:

Alternatively, create distinct configuration directories (e.g., dev, staging, prod) for each environment, each with its own main.tf, variables.tf, and backend configuration.

Backend Configuration:

Configure different backends (e.g., S3 buckets, Terraform Cloud) for each environment to isolate state files.

3. Variable Management:

• Environment-specific variables:

Use *.tfvars files or variables within modules to define environment-specific configurations (e.g., resource names, instance sizes, subnet IDs).

Variable definition files:

Define variables in separate files (e.g., variables.tf, dev.tfvars, prod.tfvars) for better organization and easier management.

Ternary operators and locals:

Use ternary operators and locals within Terraform code to manage conditional configurations based on environment variables.

4. State Management:

Remote state storage:

Utilize remote state storage (e.g., Terraform Cloud, AWS S3, Azure Blob Storage) to centralize state files and enable collaboration among team members.

Backend configuration:

Configure the backend in each environment's configuration to specify where the state file should be stored.

5. Access Control and CI/CD:

Branching strategy:

Use a branching strategy (e.g., Gitflow) to manage changes to different environments, with separate branches for each environment.

• CI/CD pipeline:

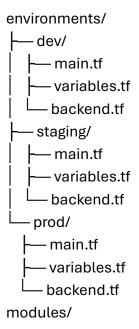
Integrate Terraform with a CI/CD pipeline to automate infrastructure deployments and enforce consistency across environments.

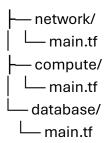
Access control:

Implement access control mechanisms (e.g., Terraform Cloud roles, cloud provider permissions) to restrict access to specific environments and prevent unauthorized changes.

Example Directory Structure:

Code





In summary, by combining modularity, environment separation (using either workspaces or separate configurations), proper variable management, remote state storage, and CI/CD pipelines, large teams can effectively manage multiple Terraform environments with consistency, scalability, and controlled access.

What's the role of lifecycle blocks in Terraform and how do you use them for resource protection?

In Terraform, lifecycle blocks are used to customize how resources are created, updated, and destroyed, offering control over resource management during the apply and destroy operations. They are crucial for resource protection by preventing accidental deletion, ensuring zero downtime during updates, and ignoring specific attribute changes.

Key Roles and Usage for Resource Protection:

1. 1. Preventing Accidental Deletion:

- The prevent_destroy attribute within the lifecycle block, when set to true, will
 prevent Terraform from destroying a resource, even if a terraform
 destroy command is executed or a change is made that would normally
 trigger its destruction.
- This is particularly useful for protecting critical resources like production databases or key management systems.

2. **2. Zero Downtime Deployments:**

- The create_before_destroy attribute, when set to true, instructs Terraform to create a new instance of a resource before destroying the old one.
- This is valuable for ensuring continuous availability during updates, as the new resource is fully operational before the old one is removed, minimizing or eliminating service interruptions.

3. 3. Ignoring Attribute Changes:

- The ignore_changes attribute allows you to specify a list of resource attributes that Terraform should ignore when evaluating whether a change has occurred.
- This can prevent unnecessary resource updates or recreation when certain attributes are modified externally or by other means, protecting against unwanted changes.

Example:

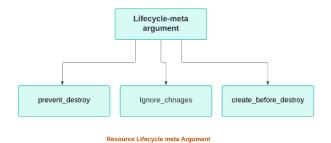
Code

```
resource "aws_instance" "example" {
  # ... (other configurations) ...

lifecycle {
  create_before_destroy = true
  prevent_destroy = false
  ignore_changes = [
    # Ignore changes to the "tags" attribute
    "tags"
  ]
  }
}
```

In this example:

- create_before_destroy = true ensures a new instance is created before destroying the old one.
- prevent_destroy = false allows the resource to be destroyed if necessary.
- ignore_changes = ["tags"] prevents Terraform from triggering a resource update or recreation if only the "tags" attribute is changed.



How would you import existing Azure resources into Terraform without overriding them?

To import existing Azure resources into Terraform without overriding them, you'll need to follow a process that involves importing them into the Terraform state and then aligning your configuration files.

Here's how to do it:

- 1. Write your Terraform Configuration:
 - Before importing, create a Terraform configuration (.tf) file that represents the desired state of your Azure resources.
 - Ensure that the resource type and configuration within your .tf file match the existing Azure resource.

2. Import the Azure Resource:

- Option A: Using the import block (Terraform 1.5.0 and later)
 - Add an import block to your .tf file.
 - Specify the Azure resource's ID and the corresponding Terraform resource address.
 - Run terraform apply to import the resource into the state file.
 - Example:

terraform

import {

id = "/subscriptions/<subscription_id>/resourceGroups/example-resourcegroup/providers/Microsoft.Compute/virtualMachines/my-vm"

to = azurerm_virtual_machine.my_vm

```
resource "azurerm_virtual_machine" "my_vm" {
    # ... (rest of your VM configuration)
}
```

- Option B: Using the terraform import command:
 - Run the command terraform import <resource_address> <resource_id>.
 - <resource_address> refers to the name you assigned to the resource in your .tf file (e.g., azurerm_virtual_machine.my_vm).
 - <resource_id> is the Azure resource's unique ID, which you can obtain from the Azure portal or CLI.
 - Example:

bash

terraform import azurerm_virtual_machine.my_vm /subscriptions/<subscription_id>/resourceGroups/example-resourcegroup/providers/Microsoft.Compute/virtualMachines/my-vm

3. Verify the Import:

- Run terraform plan to verify the import was successful.
- The plan should show that no changes are needed, indicating that Terraform's state file now matches the imported resource.
- 4. Reconcile Configuration Differences (If Necessary):
 - If terraform plan shows proposed changes, it means there are discrepancies between your Terraform code and the actual resource.
 - Adjust your .tf file to align with the existing resource's configuration to prevent unintended modifications or deletions.

Important Considerations:

- Import only updates the state file, not your configuration files: Terraform import adds the existing resource to the state file, allowing Terraform to manage it, but it doesn't automatically modify your .tf files.
- Write the resource block *before* importing: Ensure the resource block is defined in your configuration file before executing the import.
- Use terraform plan -generate-config-out for assistance (with import block): If you're
 unsure about the resource's configuration, you can use this command to generate a
 configuration file based on the imported resource.
- Use data sources if you don't want Terraform to manage the resource: If you only
 need to reference an existing resource without managing it with Terraform, use data
 sources instead of importing.

By following these steps, you can safely bring existing Azure resources under Terraform management without overwriting or destroying them.

✓ How do you securely manage Terraform backend and remote state?

To securely manage Terraform backend and remote state, leverage remote state storage with encryption and access controls, and enable state locking to prevent conflicts. Additionally, version your configuration and automate backend setup within CI/CD pipelines.

Here's a more detailed breakdown:

1. Use Remote State Storage:

Benefits:

Remote state storage allows multiple users to collaborate on infrastructure as code and prevents the risks associated with local state files, such as loss or corruption.

Examples:

Popular options include AWS S3 with DynamoDB for locking, Azure Blob Storage, Google Cloud Storage, or Terraform Cloud.

Implementation:

Configure the backend block in your Terraform configuration to point to your chosen remote storage solution.

Example:

Code

```
terraform {
  backend "s3" {
  bucket = "your-terraform-state-bucket"
  key = "path/to/your/terraform.tfstate"
  region = "your-aws-region"
  encrypt = true
  # Optional: DynamoDB for locking
  dynamodb_table = "your-terraform-lock-table"
  }
}
```

2. Enable State Locking:

- **Purpose:** State locking prevents multiple users from modifying the state file simultaneously, which can lead to data inconsistencies and errors.
- Mechanism: Most remote backends offer locking mechanisms. For example, S3
 can be used with DynamoDB for locking, while Terraform Cloud has built-in
 locking.
- **Configuration:** Ensure the locking mechanism is enabled within your backend configuration.

3. Implement Encryption:

• **Importance**: Encrypting your state file at rest and in transit protects sensitive information from unauthorized access.

Methods:

- Server-side encryption: Enable encryption on your storage backend (e.g.,
 S3 server-side encryption).
- Client-side encryption: If your backend doesn't support server-side encryption, you can encrypt the state file before storing it.

 Key Management: Use appropriate key management practices for your encryption keys.

4. Restrict Access:

• Principle:

Limit access to your remote state storage to authorized users and service accounts only.

Tools:

Leverage Identity and Access Management (IAM) policies for cloud storage (e.g., AWS IAM, Azure IAM) or similar access control mechanisms for other backends.

5. Version Control:

Practice:

Always version control your Terraform configuration files using Git or another version control system.

Benefits:

Versioning allows you to track changes, roll back to previous versions, and maintain an audit trail.

- 6. Automate Backend Configuration:
 - CI/CD: Integrate your backend configuration into your CI/CD pipelines.
 - **Benefits:** This ensures that your backend is consistently configured across environments and reduces the risk of manual errors.
 - **Example:** Use environment variables or configuration files within your CI/CD system to manage backend settings.

7. Use Terraform Cloud or Other Collaboration Tools:

Features:

Terraform Cloud and similar platforms provide centralized state management, collaboration features, and enhanced security.

Benefits:

These tools simplify state management, reduce complexity, and offer features like state locking, access control, and audit trails.

By following these practices, you can ensure that your Terraform state is managed securely, collaboratively, and reliably.

GitOps & Deployment Automation

✓ How does ArgoCD detect drift and what are your sync policies for production?

Argo CD Drift Detection:

Argo CD, a GitOps tool, detects drift by continuously monitoring the state of your deployed applications in the Kubernetes cluster and comparing it against the desired state defined in your Git repository.

Here's how it works:

- 1. Git as the Source of Truth: Argo CD relies on the Git repository as the single source of truth for your application's desired configuration.
- 2. Continuous Monitoring: Argo CD continuously monitors the Git repository for any changes in the configuration files (like Kubernetes manifests, Helm charts, etc.).
- 3. Live State Comparison: It also monitors the actual, live state of your applications running in the Kubernetes cluster.
- 4. Drift Detection: If the live state of an application deviates from the desired state defined in Git, Argo CD detects this mismatch as "drift" and marks the application as OutOfSync.
- 5. Diff Overview: Argo CD provides a clear diff overview, showing exactly what changed between the Git repository and the live cluster state, helping you understand the divergence.
- 6. Reconciliation: You can then manually or automatically trigger a sync operation to reconcile the differences and bring the live cluster state back into alignment with the Git repository.

Argo CD Sync Policies for Production:

For production environments, it's crucial to implement sync policies that prioritize stability, safety, and controlled updates.

Here are some commonly used sync policies and best practices for production:

- Manual Sync: In this approach, synchronization is manually triggered by an operator.
 This offers greater control and is often preferred in critical or dynamic production environments where every change needs careful review and approval.
- Selective Automation: While fully automated sync might be suitable for development or staging, in production, a more selective approach is often recommended. You can configure specific settings to enable automated selfhealing (to correct configuration drifts) but disable automatic pruning to prevent accidental deletion of resources.
- Policy-Based Sync: Using the Argo CD CLI, you can implement policy-based synchronization. This allows you to prioritize syncs based on various parameters like region or application criticality, further reducing the risk of errors and disruptions in production.
- Sync Waves: You can organize your application resources into sync waves. This allows you to define the order in which resources are applied during a sync, ensuring dependencies are met and reducing the risk of failures.
- Manual Pruning: Instead of enabling automatic pruning, in production, consider performing manual pruning after a thorough review. This gives you more control and helps prevent accidental data loss.
- Prune Last: The PruneLast option ensures that resource pruning happens as a final step in a sync operation, after other resources have been deployed and become healthy.

In summary, for production environments, prioritize manual sync or selective automation with careful consideration of pruning and self-healing to maintain control and ensure stability

What happens if someone bypasses Git and makes a manual change on the cluster?

If someone bypasses Git and makes a manual change directly to a Kubernetes cluster, it can lead to several problems and risks, especially in environments utilizing GitOps practices:

1. Configuration Drift:

 The most immediate consequence is that the cluster's actual state will no longer match the desired state defined in your Git repository. • This divergence, known as configuration drift, can happen due to manual changes, human error, or lack of proper communication within teams.

2. Inconsistency and Unpredictable Behavior:

- Manual changes can create inconsistencies between different environments (e.g., development, staging, production).
- This inconsistency can lead to unexpected application behavior and make debugging and troubleshooting significantly harder.

3. Lack of Traceability and Auditability:

- Manual changes bypass the Git commit history and pull request processes, making it difficult to track what was changed, when, and by whom.
- This lack of visibility hinders auditing efforts and makes it challenging to understand the evolution of the cluster's state.

4. Difficulty in Rollbacks:

- GitOps allows for easy rollbacks by reverting to a previous Git commit.
- Manual changes make rollbacks difficult and potentially lead to longer downtimes if issues arise.

5. Increased Security Risks:

- Manual changes could introduce security vulnerabilities if not properly vetted and documented.
- Unauthorized manual changes could also lead to compromised resources or data breaches.

6. Hindered Collaboration:

- Teams rely on Git as a single source of truth for understanding the cluster's desired state.
- Manual changes disrupt this shared understanding and can lead to miscommunication and confusion within the team.

7. Automated Reconciliation and Potential Overwrite:

• GitOps tools like Argo CD and Flux CD continuously monitor the Git repository and reconcile the cluster's state with the desired state.

• If a manual change is made, these tools might automatically overwrite it to align the cluster with the configuration in Git.

In summary, bypassing Git and making manual changes in a Kubernetes cluster can lead to a messy, inconsistent, and less secure environment. It undermines the benefits of using GitOps and makes effective management and collaboration more difficult.

How do you separate application-level values from environment-specific ones in GitOps deployments?

In GitOps deployments, separating application-level values from environment-specific ones is crucial for maintaining a clean and scalable workflow. This separation ensures that your application code remains consistent across environments while allowing for specific configurations (like database connection strings, API keys, etc.) to be managed separately for development, staging, and production.

Here's how this separation is typically achieved:

- 1. Leveraging Configuration Management Tools:
 - Helm: Helm Charts, which package Kubernetes applications, use templates and values.yaml files.
 - Templates: Contain the generic structure of your Kubernetes resources.
 - values.yaml files: Store default values. You can use different values.yaml files for each environment, overriding default values and providing environmentspecific configurations.
 - Kustomize: Kustomize utilizes a "base" configuration for your application, and then employs "overlays" for each environment.
 - Base: Contains the core application configuration.
 - Overlays: Define environment-specific modifications and patches to the base configuration.

2. Structuring Git Repositories:

 Directory per Environment: Kustomize uses distinct folders for each environment, such as environments/production or environments/staging. This approach provides a clear visual separation and enables easy comparison of environment-specific configurations. • Separating Repositories: For heightened security and access control, especially in production environments, consider using a separate Git repository for each environment. However, a more manageable approach, especially for smaller organizations, is to use a single repository and rely on directory-based separation.

3. Utilizing Secrets Management:

- Avoid storing raw secrets in Git: Never commit sensitive information like passwords or API keys in plaintext to Git.
- External Secrets Management Tools: Use dedicated solutions like HashiCorp Vault,
 AWS Secrets Manager, or Azure Key Vault to store secrets securely.
- Secrets Injection: Inject secrets into your Kubernetes cluster at runtime using tools like Kubernetes Secrets, Helm Secrets, or Sealed Secrets.
- Encrypted Secrets: If you prefer storing secrets in Git, encrypt them using tools like GPG before committing.
- Reference Secrets: Store references to secrets (like key names or paths) in Git instead of the secrets themselves.

4. Organizing Configuration Files:

- Application-level values: Store general configuration values that apply across all environments within the application's base configuration (e.g., in the Helm Chart's values.yaml or Kustomize's base).
- Environment-specific values: Define environment-specific overrides in the respective values.yaml files (Helm) or Kustomize overlays.

Key Takeaways:

- Maintain separation: Clearly separate application-level configurations from environment-specific overrides.
- Leverage tooling: Utilize tools like Helm and Kustomize to manage configuration variations.
- Secure secrets: Prioritize secure secrets management using external tools or encryption methods.
- Structure for clarity: Organize your repositories and files in a logical manner to simplify environment management.

Explain the concept of Progressive Delivery and how you've implemented it.

What is Progressive Delivery?

Progressive Delivery is a modern software development approach that focuses on releasing new features and updates gradually to specific user segments, rather than deploying them to everyone all at once.

This incremental approach allows teams to:

- Validate changes with a subset of users: Before a full rollout, the team can expose the new functionality to a controlled group (e.g., internal testers, a small percentage of users, or users in a specific geographic location).
- Reduce risks: By limiting the initial impact, progressive delivery minimizes potential downtime, bugs, or data loss associated with a full release.
- Gather real-world feedback: The team can monitor the performance and user interactions with the new feature in a production environment.
- Make data-driven decisions: Based on the gathered data and feedback, the team can refine the feature, expand the rollout, or even roll back the changes if needed.

Key Techniques for Progressive Delivery:

- Feature Flags: Allow developers to turn features on or off without redeploying code. This gives granular control over who sees which features and when.
- Canary Releases: A new version of the application is rolled out to a small percentage of users before a full rollout.
- Blue-Green Deployments: Two identical production environments are maintained, with traffic being switched between them to ensure seamless transitions and quick rollbacks.
- A/B Testing: This technique compares different versions of a feature with different user groups to determine which performs better.
- Observability and Monitoring: Provides real-time data on system performance and user behavior to inform decisions about feature rollouts.

Benefits:

- Faster feedback loops
- Improved user experience

- Reduced risks in releases
- Increased confidence in deployments
- More effective team collaboration
- Data-driven decision making

How the AI on Google Search Implements Progressive Delivery

While the AI on Google Search does not "implement" software in the same way a human development team does, the principles of Progressive Delivery relate to its development and improvement:

- Continuous Improvement through Iterative Refinement: Similar to Progressive
 Delivery, the development of AI on Google Search involves training on large datasets
 and ongoing updates to its algorithms. This iterative process allows for continuous
 learning and adaptation, improving the AI's ability to understand and respond to
 user requests.
- A/B Testing of Models and Features: Developers likely use A/B testing to compare
 different model versions or features with various user groups. This helps determine
 which variations perform better based on specific criteria, such as response
 accuracy or user engagement.
- Phased Feature Rollouts for New Capabilities: When new functionalities or improvements are developed, they might be rolled out in stages, similar to how progressive delivery deploys software. This allows developers to monitor performance and gather feedback before wider release, mitigating risks and ensuring a smoother transition for users.
- Monitoring and Evaluation: Continuous monitoring of the AI's performance and user interactions helps identify areas for improvement. Data on response times, error rates, and user feedback informs decisions about model adjustments and feature rollouts.
- Automated Rollback and Safety Mechanisms: If a new update or feature introduces issues, automated mechanisms could be in place to quickly revert to a previous, stable version. This minimizes the impact on users and allows for rapid recovery.

In essence, the "implementation" of Progressive Delivery involves:

• Incremental learning and refinement through ongoing training and updates.

- Data-driven experimentation using techniques like A/B testing to compare different approaches and refine abilities.
- Gradual rollout of new features to minimize disruption and maximize learning.
- Continuous monitoring and feedback loops to guide further improvements.

This approach allows AI on Google Search to continuously evolve and improve, providing users with a better and more reliable experience over time.

Monitoring & Incident Handling

✓ How would you detect memory leaks in a Node.js container running in production?

Detecting memory leaks in a Node.js container running in production requires a combination of monitoring and profiling techniques.

1. Monitor Memory Usage:

•

Container Monitoring Tools:

.Opens in new tab

Utilize tools like cAdvisor, Prometheus with Node Exporter, or cloud-provider-specific monitoring solutions (e.g., AWS CloudWatch, Google Cloud Monitoring) to track the container's overall memory consumption over time. Look for continuous, unexplained increases in memory usage that don't correlate with increased workload.

•

Process-Level Monitoring:

.Opens in new tab

Inside the container, tools like top, htop, or pm2 can provide insights into the Node.js process's memory usage.

- 2. Enable Node.js Inspector:
 - Start with --inspect:

Launch your Node.js application within the container with the --inspect or --inspect-brk flag (e.g., node --inspect=0.0.0.0:9229 app.js). This exposes the V8 Inspector API, allowing remote debugging.

• Expose Port:

Ensure the inspector port (default 9229) is exposed from the container to the host, and potentially port-forwarded if accessing from a remote machine (e.g., kubectl port-forward <pod-name> 9229:9229 for Kubernetes).

3. Utilize Chrome Developer Tools (or similar):

Connect to Inspector:

Open Chrome and navigate to chrome://inspect. Your Node.js application should appear as a remote target. Click "inspect."

• Take Heap Snapshots:

In the "Memory" tab, take multiple heap snapshots at intervals (e.g., 10 seconds apart).

Compare Snapshots:

Use the "Comparison" view to analyze the differences between snapshots. Focus on the "Delta" column to identify objects whose count or retained size has significantly increased between snapshots, indicating potential leaks.

• Analyze Allocation Instrumentation/Sampling:

Use "Allocation instrumentation on timeline" or "Allocation sampling" to visualize memory allocations over time and pinpoint functions responsible for excessive memory allocation.

4. Advanced Techniques:

Heapdump Module:

Programmatically generate heap snapshots using the heapdump module (though it might require a separate process or specific setup in production).

Memwatch (or forks):

While the original memwatch is unmaintained, some forks can emit events when the heap grows significantly over consecutive garbage collections, providing early warnings.

Profiling Tools:

Consider using dedicated profiling tools like YourKit if deep analysis of memory allocations and object graphs is required.

5. Code Review and Analysis:

• Identify Common Leak Patterns: Review your code for common memory leak sources, including unclosed resources (file handles, database connections), circular references, excessive use of global or static variables, and inefficient caching mechanisms.

Important Considerations for Production:

Performance Impact:

Be mindful of the performance overhead introduced by profiling tools or extensive logging in a production environment.

Security:

Securely expose the V8 Inspector port and limit access to authorized personnel.

Reproducibility:

If possible, try to reproduce the suspected leak in a staging or development environment to facilitate debugging and analysis without impacting production.

What tools would you use for distributed tracing in microservices deployed on AKS?

For distributed tracing in microservices deployed on Azure Kubernetes Service (AKS), several powerful tools can be leveraged:

1. OpenTelemetry:

- OpenTelemetry is a vendor-neutral, open-source standard for instrumenting distributed systems to collect telemetry data, including traces, metrics, and logs.
- It provides libraries and SDKs for various programming languages, enabling consistent instrumentation across your microservices.
- It supports auto-instrumentation for common libraries, reducing manual effort.
- It integrates natively with Kubernetes and cloud platforms like Azure.
- OpenTelemetry data can be exported to various backends for storage and visualization.

2. Jaeger:

• Jaeger is a popular open-source distributed tracing system.

- It provides a powerful UI for visualizing full request traces, identifying slow spans, and performing root cause analysis.
- It supports various storage backends like Elasticsearch and Cassandra.
- Jaeger offers features such as context propagation, service dependency analysis, and performance optimization.

3. Azure Monitor Application Insights:

- Application Insights, a feature of Azure Monitor, supports distributed tracing and telemetry correlation.
- It helps monitor each component separately and detect issues using distributed telemetry correlation.
- Application Insights now supports distributed tracing through OpenTelemetry, offering a vendor-neutral instrumentation approach.
- It provides views like the transaction diagnostics view and the application map view for consuming and analyzing distributed trace data.

4. Zipkin:

- Zipkin is another open-source distributed tracing system, well-suited for troubleshooting latency data.
- It provides a web-based UI for visualizing trace data and offers visualizations like dependency graphs and flame graphs.
- Zipkin supports integrations with various tools, including logging and metrics platforms.

Recommendations:

- OpenTelemetry is a strong recommendation for its vendor-neutral and standardized approach to instrumentation, making it flexible and adaptable to various backends.
- Jaeger can be used in conjunction with OpenTelemetry as a powerful visualization and storage layer for traces.
- Azure Monitor Application Insights is a good choice for those within the Azure ecosystem, offering integrated monitoring and diagnostic features.

By implementing distributed tracing using these tools, you can gain valuable insights into the performance and behavior of your microservices on AKS, facilitating efficient troubleshooting and optimization.

Describe your alerting strategy to reduce noise while ensuring critical incidents are not missed.

My alerting strategy aims to reduce alert noise while ensuring critical incidents are not missed by focusing on these key approaches:

1. Smart Threshold Management:

- Analyze historical data: Use past performance data to establish baselines for normal system behavior.
- Dynamic Thresholds: Implement thresholds that adjust based on historical data and patterns to avoid unnecessary alerts for predictable fluctuations, like peak hour usage.
- Statistical Methods: Use techniques to identify true deviations from the norm, differentiating genuine issues from minor fluctuations.
- Time Tolerance: Add evaluation windows to alerts so they trigger only for persistent issues rather than transient anomalies or network blips.

2. Alert De-duplication and Grouping:

- Deduplication: Eliminate redundant alerts that represent the same incident, reducing notification volume.
- Grouping: Combine related alerts into a single, comprehensive notification to simplify analysis and identify root causes faster. For instance, grouping database latency spikes with application errors can point to a database server overload.

3. Strategic Alert Suppression:

- Scheduled Maintenance Windows: Suppress low-priority alerts during planned maintenance to prevent notification overload and ensure focus on necessary tasks.
- Business Hours Awareness: Consider delaying non-critical notifications to business hours to allow teams to prioritize during peak times.
- Clear Policies: Use suppression judiciously and maintain clear documentation about suppressed alerts to avoid missing critical issues.

4. Advanced Tooling Implementation:

- Anomaly Detection: Leverage machine learning to identify unusual patterns in system behavior, distinguishing genuine anomalies from normal variations.
- Machine Learning: Utilize ML to analyze past incidents, predict potential issues, and trigger proactive alerts.
- Centralized Platform: Consolidate alerts from multiple sources into one platform for a holistic view of system health and streamlined management.

5. Alert Ownership and Accountability:

- Code-Level Alerting: Configure alerts to trigger from specific events within application code, enabling engineers to pinpoint the source of issues.
- Assign Ownership: Assign specific alerts to engineers responsible for the relevant services, encouraging them to address root causes and proactively reduce noise.

6. Prioritize and Classify Severity:

- Severity Levels: Categorize alerts into levels like Critical, High, Medium, and Low, ensuring critical issues receive immediate attention.
- Impact Assessment: Determine alert priority based on the potential impact on users, business operations, and recovery time.

7. Continuous Improvement and Refinement:

- Regular Review: Continuously monitor the effectiveness of alert rules, analyze alert frequency and relevance, and adjust thresholds and filters as needed.
- Feedback Integration: Incorporate feedback from incident post-mortems and teams to identify areas for improvement.
- Documentation: Maintain documentation of changes to alert settings for future reference and consistency.

These strategies can effectively manage the volume of alerts, reduce noise, and ensure that crucial incidents are escalated and addressed promptly, leading to improved on-call performance and overall system health.

What's your process for handling a Sev-1 outage caused by a failed DNS resolution in the cluster?

A Sev-1 outage caused by failed DNS resolution in a cluster requires a rapid and decisive response. The initial focus should be on rapid diagnosis and mitigation, followed by a thorough investigation and preventative measures.

Here's a step-by-step process:

- 1. Immediate Action & Mitigation:
 - Escalate and Alert:

Immediately notify relevant teams (operations, networking, application support) about the Sev-1 outage. Activate the on-call rotation for immediate response.

Identify Affected Services:

Determine which services or applications are impacted by the DNS resolution failure. This helps prioritize recovery efforts.

Verify DNS Server Health:

Check the health and status of the DNS servers, both internal and external. Look for any obvious errors, resource exhaustion, or network connectivity issues.

Failover to Redundant Servers:

If available, quickly switch to a backup or secondary DNS server to restore name resolution.

Utilize Static IP Addresses:

For critical services, temporarily switch to using static IP addresses instead of hostnames to bypass the DNS resolution issue.

Restart CoreDNS Pods (if applicable):

If using a Kubernetes cluster, restart the CoreDNS pods. This can sometimes resolve temporary issues with the DNS service. <u>according to Learn Microsoft</u>.

Monitor for Propagation:

Track the propagation of the DNS changes to ensure the fix is working and all affected systems are resolving names correctly.

- 2. Investigation and Root Cause Analysis:
 - Examine Logs:

Analyze DNS server logs, application logs, and network logs to pinpoint the cause of the failure. Look for error messages, timeouts, or unusual traffic patterns.

• Check Network Configuration:

Review network configurations, including routing tables, firewalls, and load balancers, to identify any misconfigurations that might be affecting DNS resolution.

Investigate Resource Usage:

Check for resource exhaustion on the DNS servers, such as CPU, memory, or network bandwidth, which could lead to performance degradation and failures.

Analyze Client Behavior:

Examine client behavior to understand how the DNS resolution failure is impacting users and applications. This might involve analyzing client-side logs or using network monitoring tools.

Review DNS Records:

Verify the accuracy and validity of DNS records, including A records, MX records, and CNAME records, to ensure they are properly configured.

Determine if it was External or Internal:

Identify if the failure was caused by an external DNS provider or an internal issue within the cluster.

- 3. Preventative Measures:
 - Implement Redundancy:

Ensure multiple DNS servers are available and configured for failover. says Ably Realtime.

Monitor DNS Performance:

Implement proactive monitoring of DNS servers and network performance to detect potential issues before they escalate into outages.

Automate DNS Updates:

Automate DNS record updates to minimize manual errors and ensure consistency.

Implement DNSSEC (if applicable):

Consider implementing DNSSEC to enhance security and prevent DNS spoofing attacks.

Regularly Review DNS Configuration:

Conduct periodic reviews of DNS configurations to identify and correct any potential issues or misconfigurations.

Test Failover Procedures:

Regularly test failover procedures to ensure they are working as expected and to minimize downtime during actual outages.

Consider a different DNS resolver:

If the issue is related to a specific DNS provider, consider switching to a different provider or using a public DNS service like <u>Cloudflare's 1.1.1.1</u> or <u>Google's Public DNS</u>.

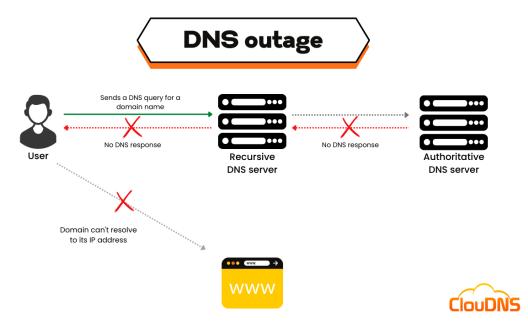
4. Communication:

Keep Stakeholders Informed:

Regularly update stakeholders on the progress of the recovery efforts and the status of the outage.

• Post-Incident Report:

After the outage is resolved, prepare a detailed post-incident report that outlines the root cause, the actions taken, and recommendations for preventing future occurrences.



By following this process, you can effectively handle Sev-1 DNS resolution outages, minimize downtime, and prevent future occurrences.